

Single-threaded CPU Approach:

The time complexity is $O(N)$.

After modifying the original naive CPU implementation to use a single loop, the time complexity is reduced to $O(N)$. For the initial version with two nested loops, it had a time complexity of $O(N^2)$.

Naive GPU Approach:

The time complexity is $O(N^2)$.

Although the algorithm itself has a time complexity of $O(N)$, utilizing N threads leads to a total time complexity of $O(N^2)$.

Recursive Doubling GPU Approach:

The time complexity is $O(N \log N)$.

The algorithm's complexity is $O(\log N)$, and with N threads, the overall time complexity becomes $O(N \log N)$.